

Introduction to Computer Vision



Matthieu Baudry - Nov. 2020 matthieu.baudry@airbus.com

Download documents on the LMS

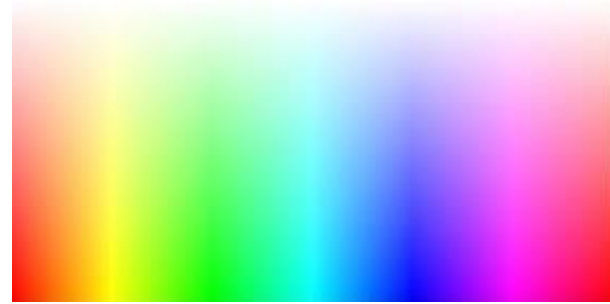
Objectives and Directives

- Understand **the basis** of Computer Vision
- Be able to
 - Modify images
 - Extract information
 - Apply Fourier Transform
 - Filter images
 - Apply deblurring
- Evaluation:
 - Report & clear code
 - Date: in 2 weeks (**13th of December**)
 - Notation:
 - Quality of redaction, form (put clear illustrations!)
 - Quality and pertinence of comments
 - “Extra work” for bonus (**test other images, filters, applications**, etc.)

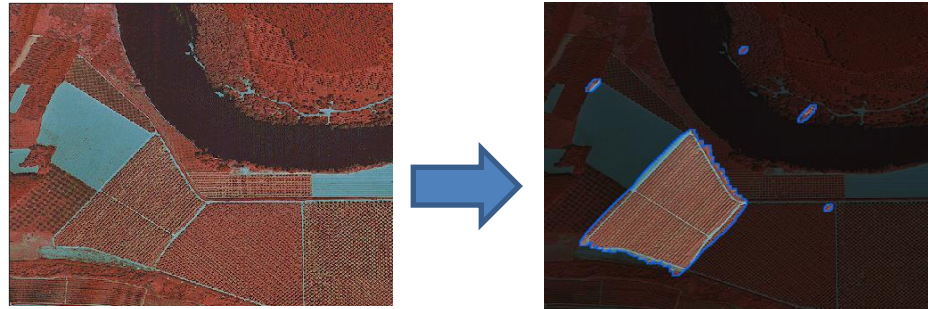
Framework

I – Basics

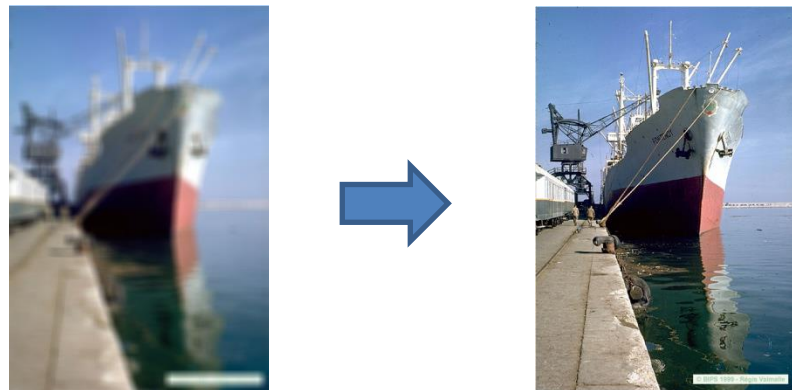
- GrayScale Image
- Coloured Image
- Histogram
- Filtering



II – Fourier Transform



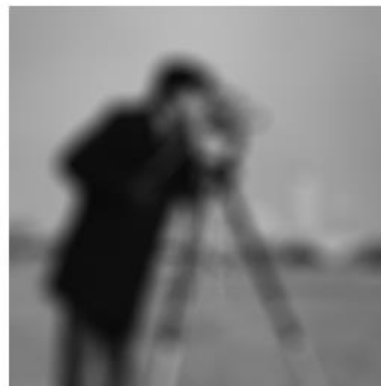
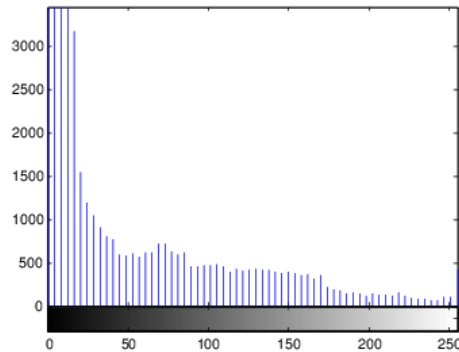
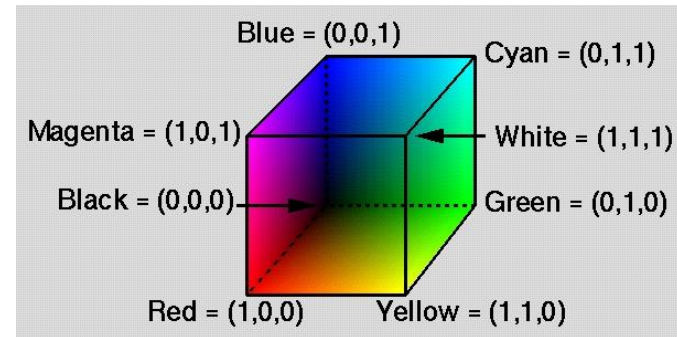
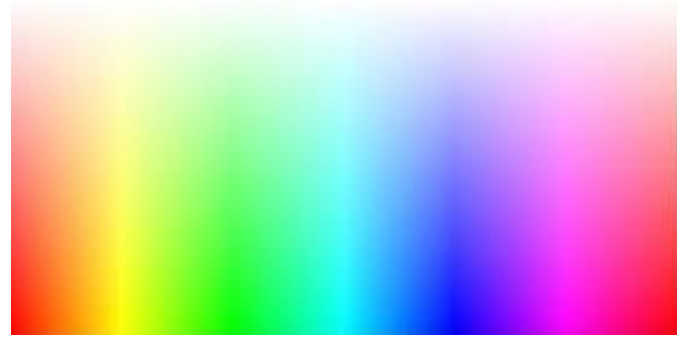
III - Deblurring



I - Basics

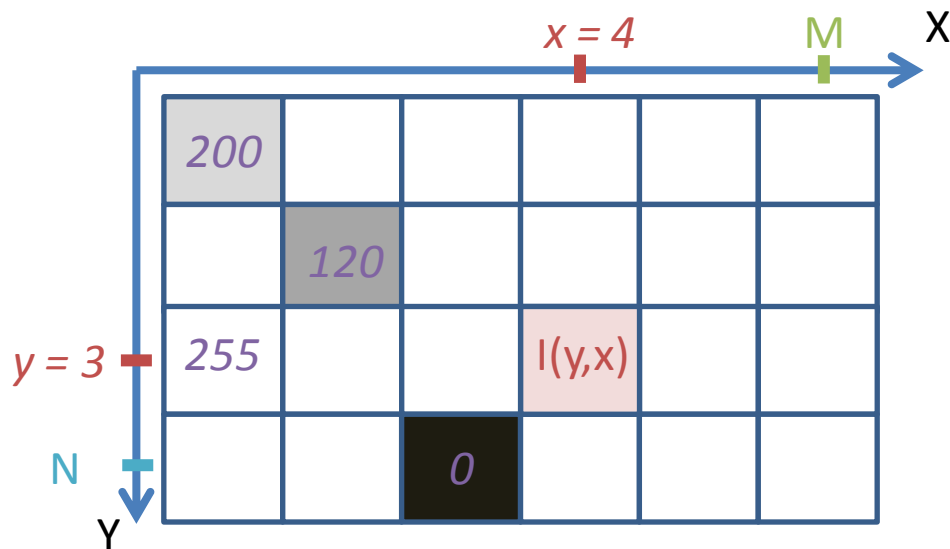
I – Basics

- GrayScale Image
- Coloured Image
- Histogram
- Filtering



I - Basics

- GrayScale Image I of size $(N \times M)$ = Matrix of $N \times M \times 1$ elements, whose values belongs to $[0:255]$ (for UINT8, $[0:1]$ for double)
- Lowest value = Dark, Highest value = white
- Ex: $N = 4$, $M = 6$



Convention depends on languages

- in C, indexes begin at 0 and end at $N-1$;
- In MatLab, indexes begin at 1 and end at N ;
- access might be $I(x,y)$ instead of $I(y,x)$;
- etc.


I - Basics

Matlab commands:

- Create an image:

- `I = zeros(N,M);` *% if $N=M$, $I = \text{zeros}(N)$*
- `I = ones(N,M);` *% Matrix filled with 1*
- `I = eye(N);` *% Diagonal matrix*
- `I = ones(N,M,3);` *% $N \times M \times 3$ matrix (for RGB or HSV images for instance)*

- Modify pixels value:

- `I(n,m) = new_value;` *% row n , column m*
- `I = 255*I;` *% Multiplies all elements of I by 255 (as a scalar)*
- `I3 = I1*I2;` *% Normal matrix operation*
- `I3 = I1.*I2;` *% Element per element operation*
- `I(n1:n2:m1:m2) = I2;` *% Put $I2$ in I .  be careful to have correct sizes*

- Load an image:

- `I = imread('path/image.png');`

- Display an image:

- `figure, imshow(I);` or `imagesc(I);` *% `imagesc` scales the color levels to values of I*
- `close all`

- Save an image:

- `imwrite(I,'MyImage.png','png');`

- Help:

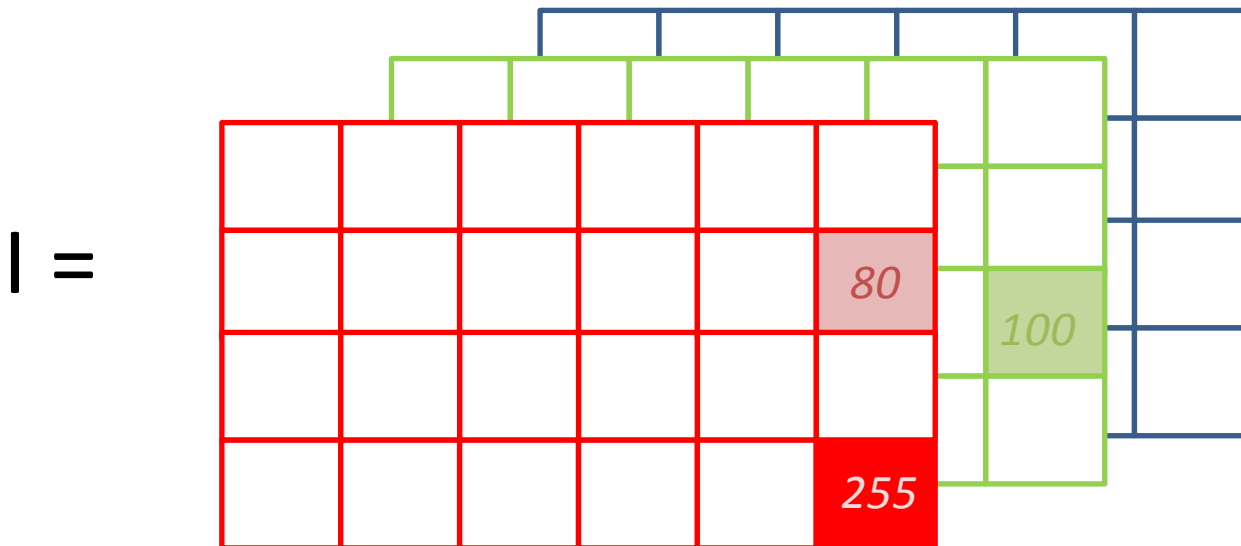
- `help imwrite;`
- `doc imwrite;`

See Code section 1 (and video 1)

I - Basics

- Coloured Image **I** of size (**N*****M**) = Matrix of **N*****M*****3** elements
- Different color-spaces exist: (RGB), (HSV), (CMYK), ...
- Ex: **RGB**

Principle: 1 Matrix for each Red, Green or Blue component

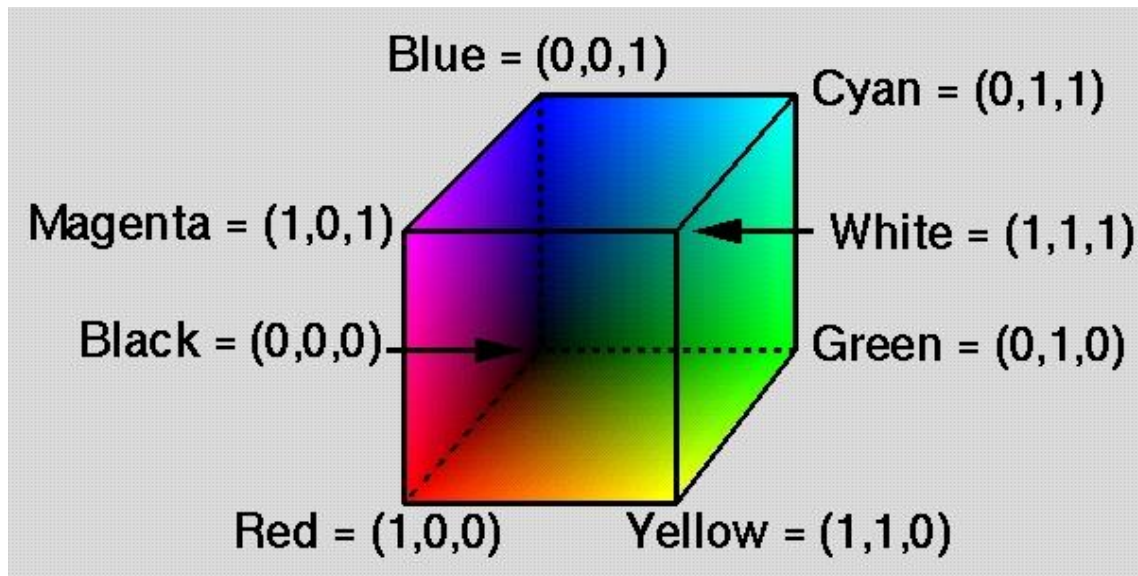


- $I_{red} = I(:, :, 1)$
- $I_{green} = I(:, :, 2)$
- $I_{blue} = I(:, :, 3)$

I - Basics

- Coloured Image **I** of size (**N*****M**) = Matrix of **N*****M*****3** elements
- Different color-spaces exist: (RGB), (HSV), (CMYK), ...
- Ex: **RGB**

Principle: 1 Matrix for each Red, Green or Blue component

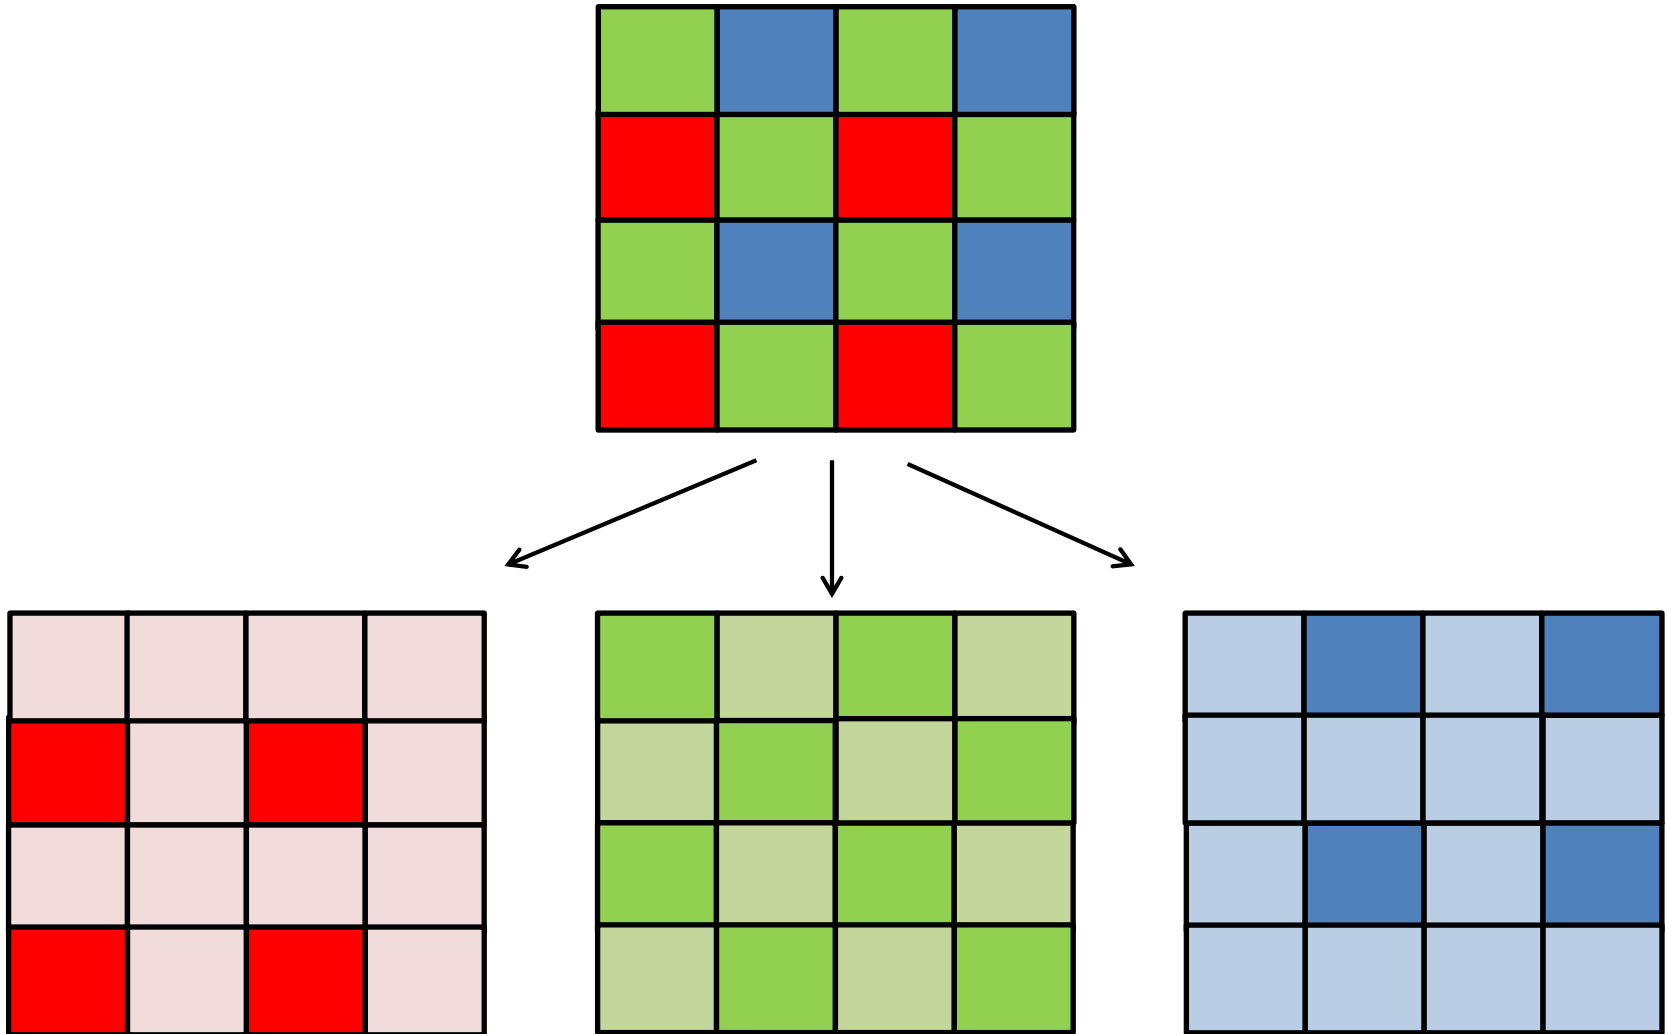


- $I_{red} = I(:, :, 1)$
- $I_{green} = I(:, :, 2)$
- $I_{blue} = I(:, :, 3)$

See Code section 2 (and video 2)

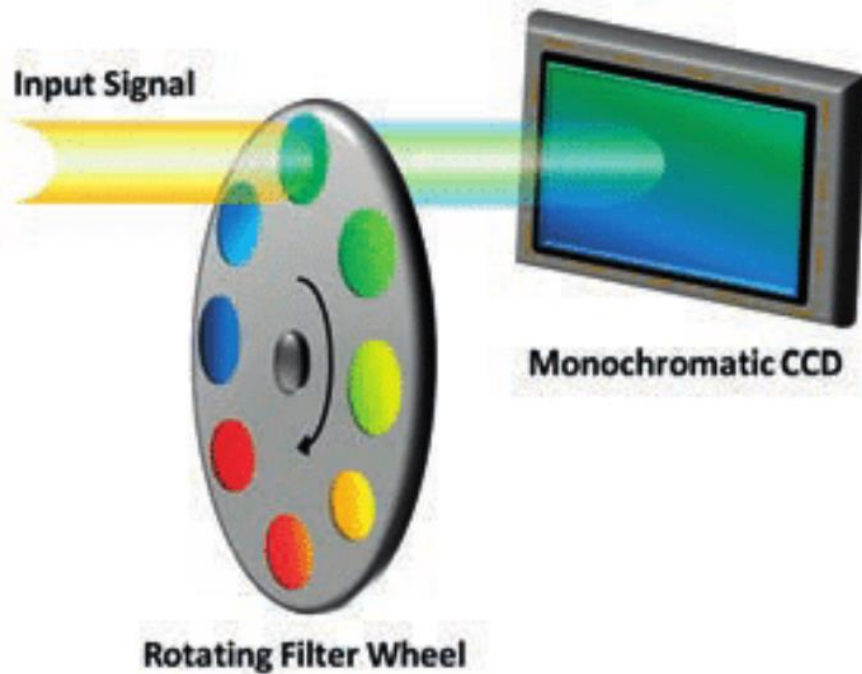
I - Basics

- Acquisition of RGB Image: Bayer matrix + interpolation



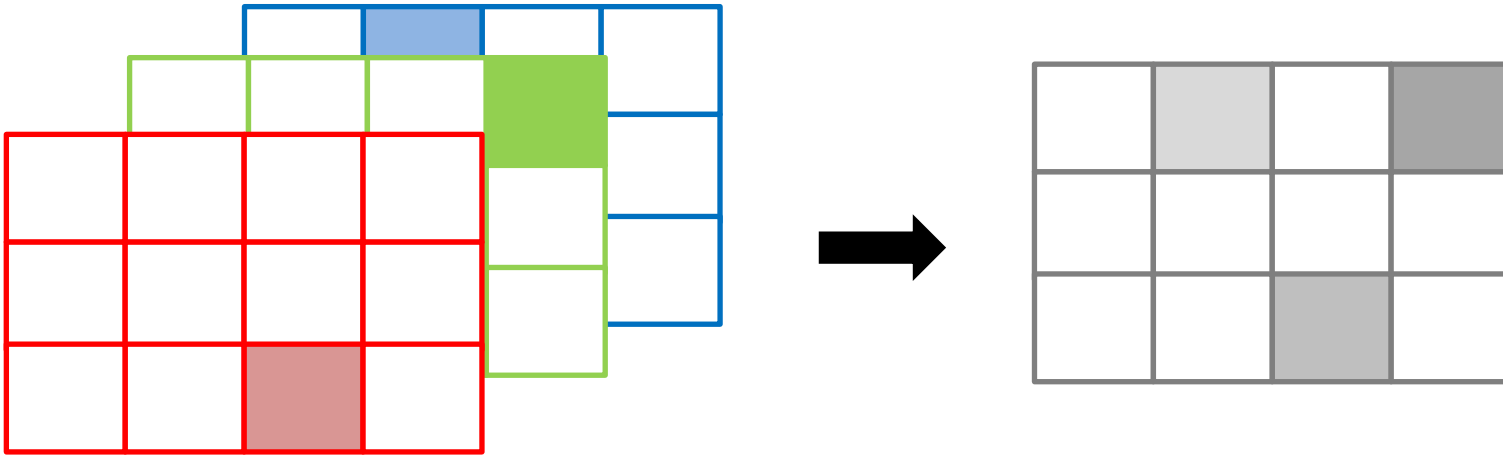
I - Basics

- Acquisition of multispectral Image: Rotating wheel



I - Basics

- RGB to Gray



$$\triangleright I_g = \alpha I_{red} + \beta I_{green} + \gamma I_{blue}$$



This transformation is not bijective (loss of information)



$$\triangleright I_g = rgb2gray(I_{RGB})$$

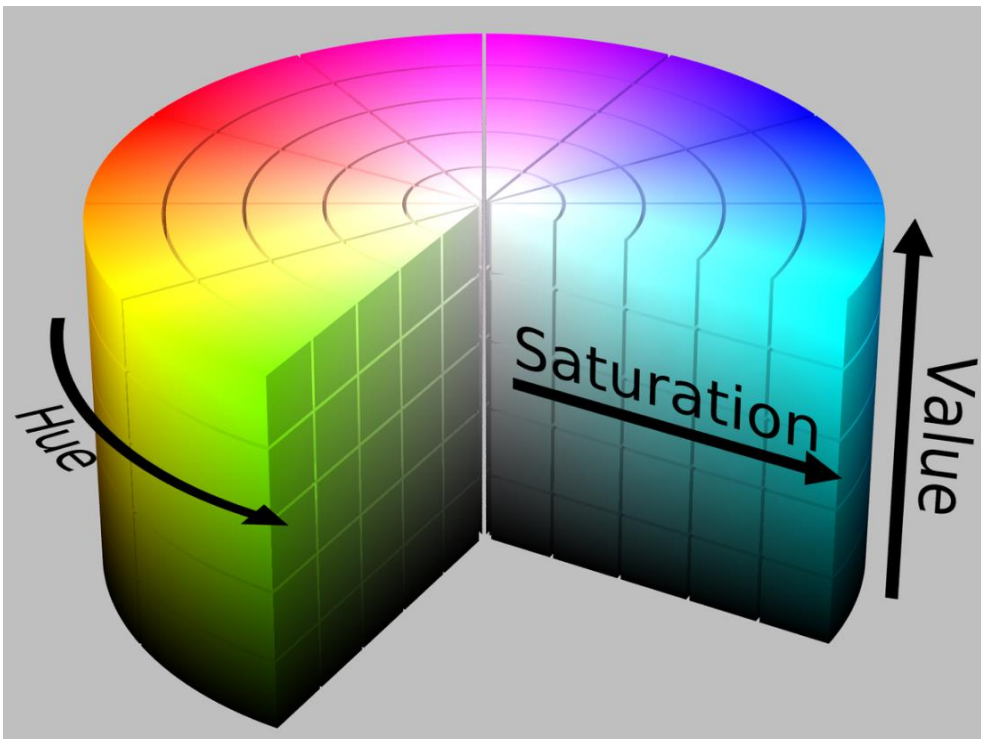
I - Basics

- HSV (Hue Saturation Value) Color space

Hue \approx Pure color (wavelength) [*teinte* in French]

Saturation \approx Intensity of coloration

Value \approx The brightness



$$\triangleright I_{HSV} = rgb2hsv(I_{RGB})$$

https://en.wikipedia.org/wiki/HSL_and_HSV

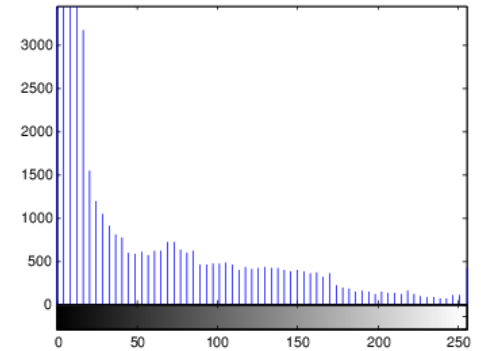
I - Basics

- Histogram

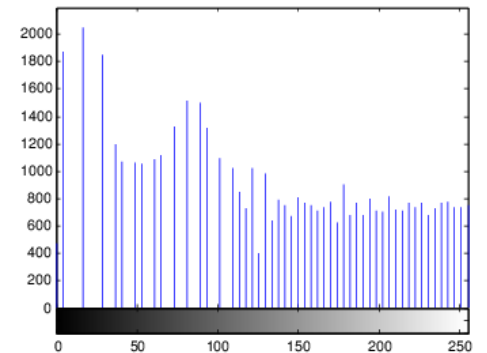
➤ $\text{Hist} = \text{imhist}(I_g)$

$\text{Hist}(i) = \text{number of pixels of } I \text{ that have a value equals to } i$

Gives information about the contrast of the image



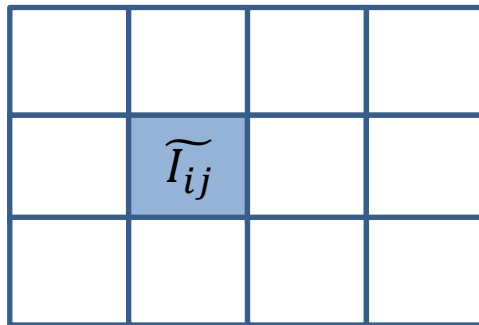
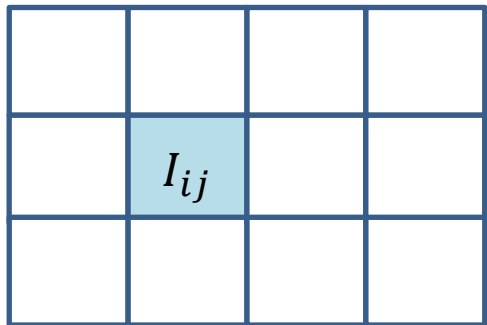
➤ $I_{eq} = \text{histeq}(I_g)$



I - Basics

- Filtering

Important part of Image Processing → modify intensity of **each pixel** in accordance with some rule (locally or globally)



Where: $\widetilde{I}_{ij} = f(I, \dots)$

Used to remove noise, sharpen edges, blur areas, etc.

Ex:



Filtering can be *linear* or *non linear*

I - Basics

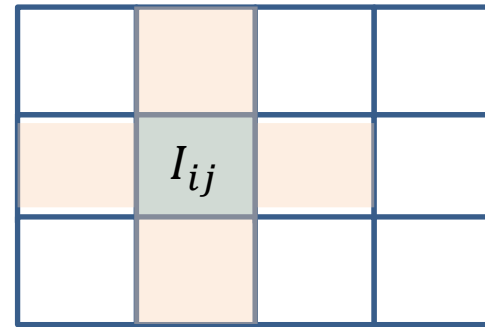
- **Non linear Filtering (1/2)**

Apply a non linear operation to the image

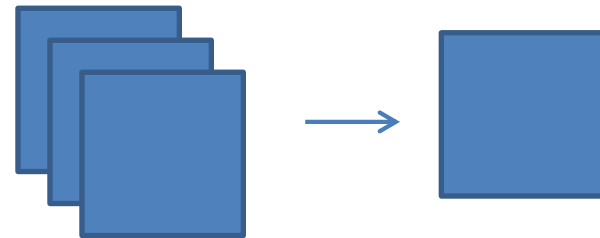
Ex:

$$I_{locmin}(i,j) = \min_{(k,l) \in H} (I(i+k, j+l))$$

$H : \{(i,j) \text{ \& 4 closest neighbours}\}$



$$I_{min}(i,j) = \min(I_1(i,j), I_2(i,j), \dots)$$



Other non linear filters:

- max,
- median,
- bilateral filter,
- etc.

I - Basics

- **Non linear Filtering (2/2)**

Ex:

Median Filtering:



https://en.wikipedia.org/wiki/Median_filter

Bilateral Filtering:



http://homepages.inf.ed.ac.uk/rbf/CVonline/LOCAL_COPIES/MANDUCHI1/Bilateral_Filtering.html

I - Basics

- **Linear Filtering (1/8)**

Apply a linear operation to the image: convolution by a kernel

Lets H be a *kernel matrix* (any matrix, size $[2n+1, 2m+1]$)

We have :

$I_{filtered} = I * H$, where $*$ is the convolution operator (2D)

Rem. Conv 1D: $(f * g)(x) = \int_{-\infty}^{+\infty} f(t)g(x - t)dt$

$$I_{filtered}(i,j) = \sum_{k=1}^{2n+1} \sum_{l=1}^{2m+1} H(k,l) . I(i + n - k + 1, j + m - l + 1)$$

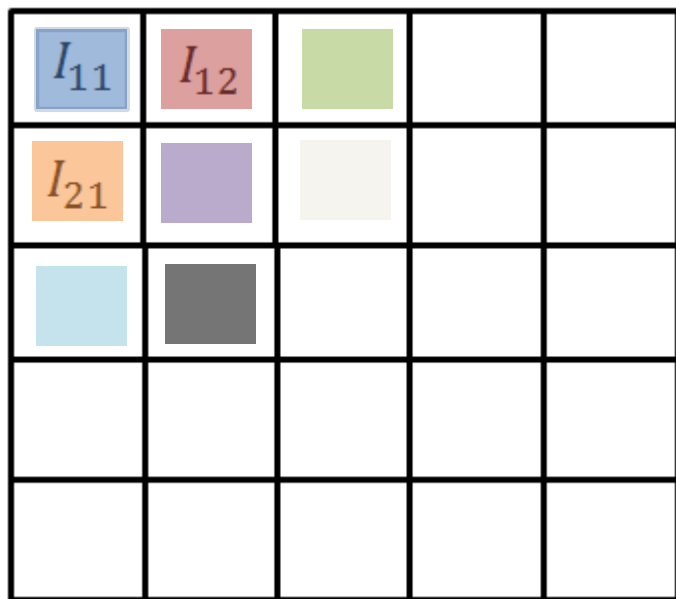
➤ $I_{filtered} = imfilter(I, H, [options])$

I - Basics

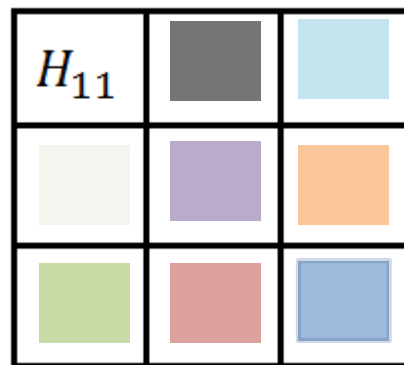
- Linear Filtering (2/8)

Ex:

$$I_{filtered}(i,j) = \sum_{k=1}^{2n+1} \sum_{l=1}^{2m+1} H(k,l) \cdot I(i+n-k+1, j+m-l+1)$$



*



$$I_{filtered}(2,2) = I_{11}H_{33} + I_{12}H_{32} + I_{13}H_{31} + I_{21}H_{23} + I_{22}H_{22} + I_{23}H_{21} + I_{31}H_{13} + I_{32}H_{12} + I_{33}H_{11}$$

≠ correlation:

$$I_{corr}(i,j) = \sum_{k=1}^{2n+1} \sum_{l=1}^{2m+1} H(k,l) \cdot I(i-n+k-1, j-m+l-1)$$

Specify option 'conv' with matlab, otherwise it performs correlation (see code sections 3&4).

I - Basics

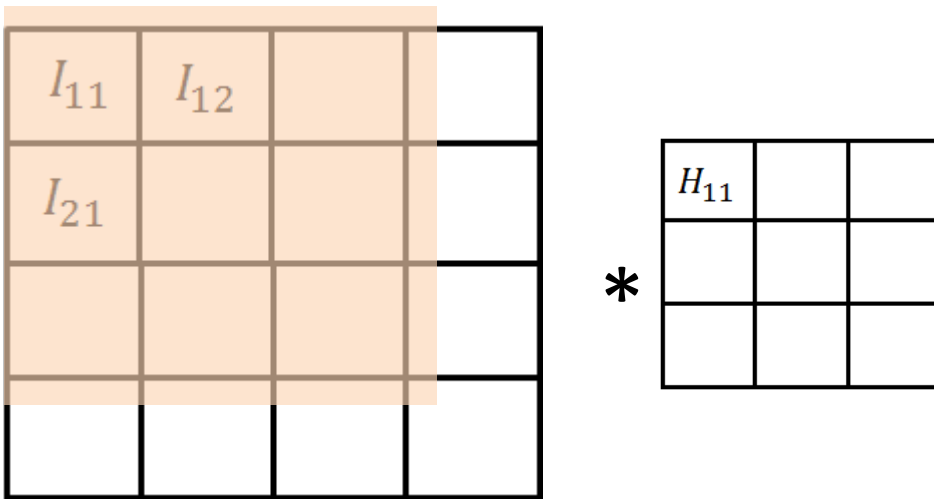
- Linear Filtering (3/8)

Ex: simple kernels

- If $H = h$ (ie scalar), then $I_{filtered} = H * I = h.I$

- If $H = \frac{1}{9} \cdot \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$, then $I_{filtered}(2,2) = \frac{1}{9} I_{11} + \frac{1}{9} I_{12} + \dots + \frac{1}{9} I_{33}$

$$= average(I_{11}:I_{33})$$

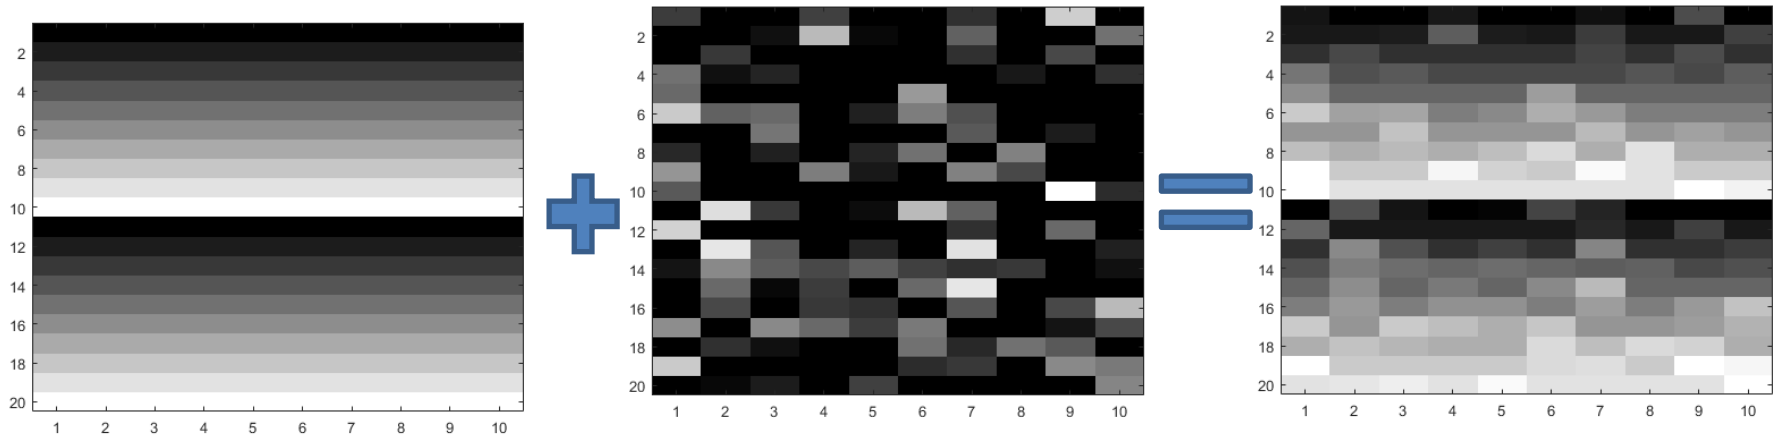


I - Basics

- Linear Filtering (4/8)

Ex: Smooth noise in image

$$I_{noisy} = I + I_{noise} \quad \text{with} \quad I_{noise}(i,j) \sim \mathcal{N}(0, \sigma^2) \quad \forall i, j$$



How to retrieve I from I_{noisy} ?

→ Use properties:

If $(x_i)_{i \in [1,m]} \sim [\mathcal{N}(0, \sigma^2)]^m$ then: $\frac{1}{m} \sum x_i \sim \mathcal{N}(0, (\sigma')^2)$ with $\sigma' = \frac{\sigma}{\sqrt{m}} \rightarrow 0$ as $m \rightarrow \infty$

I - Basics

- Linear Filtering (5/8)

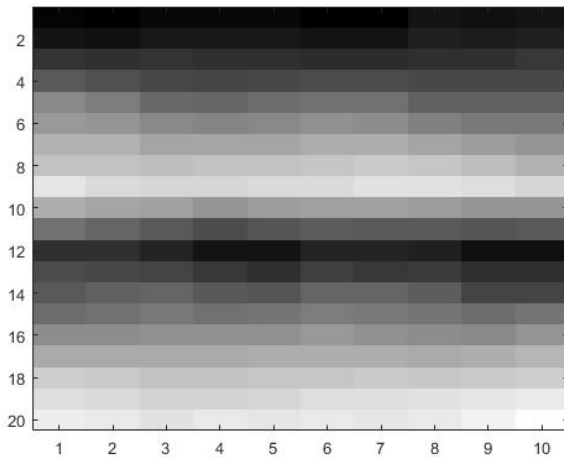
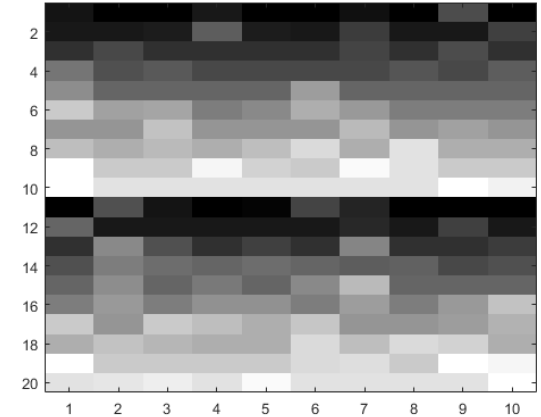
Ex: Smooth noise in image

How to retrieve I from I_{noisy} ?

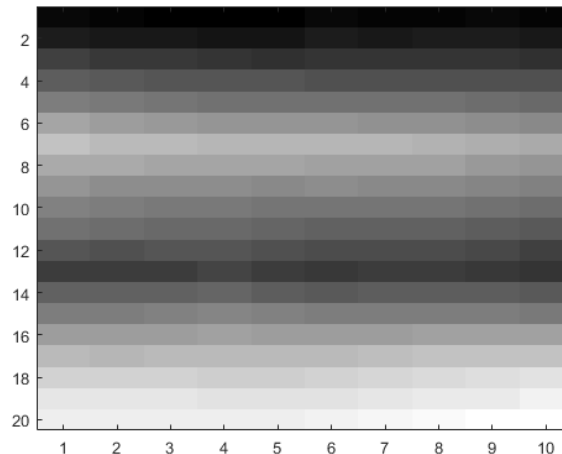
→ Average the image, for instance with a mean operator.

Concretely:

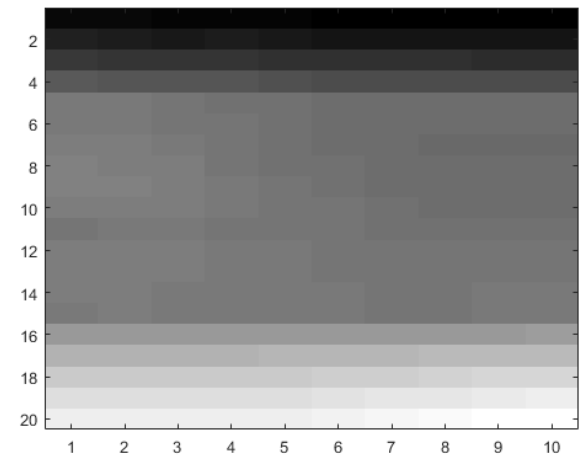
$$\text{use } H = \frac{1}{m^2} * \begin{bmatrix} 1 & \dots & 1 \\ \vdots & \ddots & \vdots \\ 1 & \dots & 1 \end{bmatrix}$$



$m = 3$



$m = 6$



$m = 10$

→ Prefer Gaussian Kernels which better preserve edges

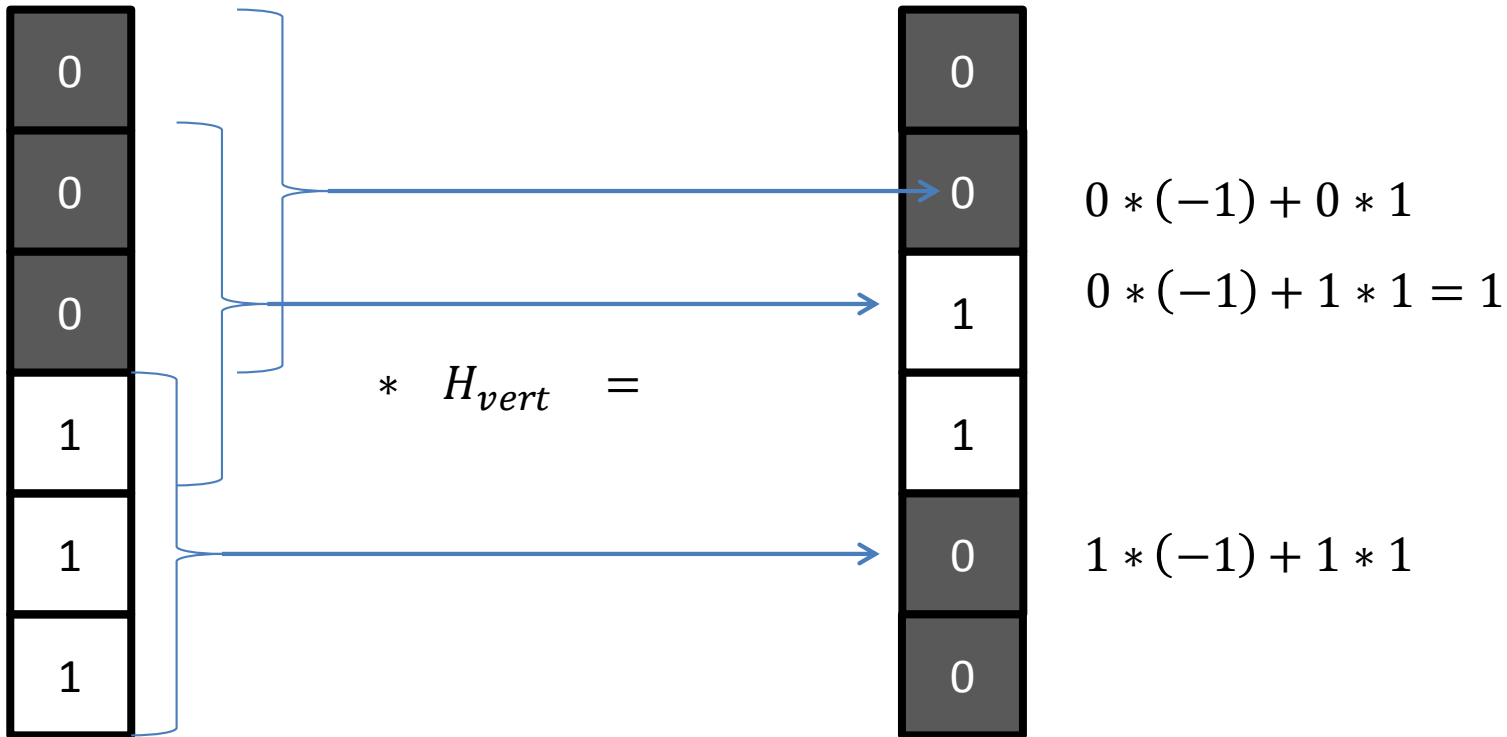
... See implemented code, section 3
(and video 3)

I - Basics

- Linear Filtering (7/8)

Ex 2 : Gradient of an image (vertical)

$$H_{vert} = \begin{bmatrix} 1 \\ 0 \\ -1 \end{bmatrix}$$



... See implemented code, section 4 (and videos 4 & 4b)

I - Basics

- Linear Filtering (7/8)

Property:

$$H_1 * (H_2 * I) = (H_1 * H_2) * I \quad (\text{here, } * \text{ operator is still convolution})$$

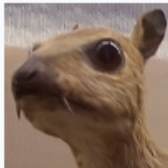



For instance:





Instead of smoothing the image **and** generate gradient (which takes time), use one Kernel that combines both operations: $H_3 = H_2 * H_2$

$$\left. \begin{array}{l} H_1 = [-1 \ 0 \ 1] \text{ (gradient)} \\ H_2 = \frac{1}{3} * \begin{pmatrix} 1 \\ 1 \\ 1 \end{pmatrix} \text{ (smooth)} \end{array} \right\} \quad H_3 = \frac{1}{3} * \begin{bmatrix} -1 & 0 & 1 \\ -1 & 0 & 1 \\ -1 & 0 & 1 \end{bmatrix} \begin{array}{l} \downarrow \text{ Noise Smoothing} \\ \longrightarrow \text{ Edge detection} \end{array}$$

I - Basics

- Linear Filtering (8/8)

Operation	Kernel ω	Image result $g(x,y)$
Identity	$\begin{bmatrix} 0 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 0 \end{bmatrix}$	
Edge detection	$\begin{bmatrix} 1 & 0 & -1 \\ 0 & 0 & 0 \\ -1 & 0 & 1 \end{bmatrix}$	
	$\begin{bmatrix} 0 & 1 & 0 \\ 1 & -4 & 1 \\ 0 & 1 & 0 \end{bmatrix}$	
	$\begin{bmatrix} -1 & -1 & -1 \\ -1 & 8 & -1 \\ -1 & -1 & -1 \end{bmatrix}$	

Sharpen	$\begin{bmatrix} 0 & -1 & 0 \\ -1 & 5 & -1 \\ 0 & -1 & 0 \end{bmatrix}$	
Box blur (normalized)	$\frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$	
Gaussian blur 3×3 (approximation)	$\frac{1}{16} \begin{bmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{bmatrix}$	
Gaussian blur 5×5 (approximation)	$\frac{1}{256} \begin{bmatrix} 1 & 4 & 6 & 4 & 1 \\ 4 & 16 & 24 & 16 & 4 \\ 6 & 24 & 36 & 24 & 6 \\ 4 & 16 & 24 & 16 & 4 \\ 1 & 4 & 6 & 4 & 1 \end{bmatrix}$	

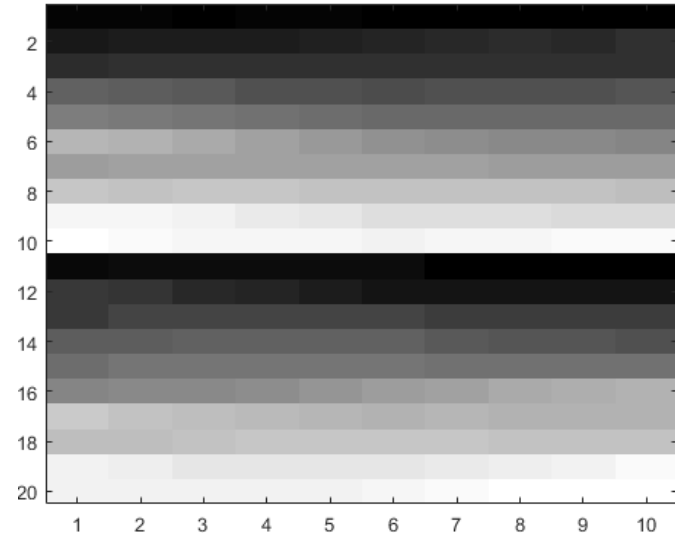
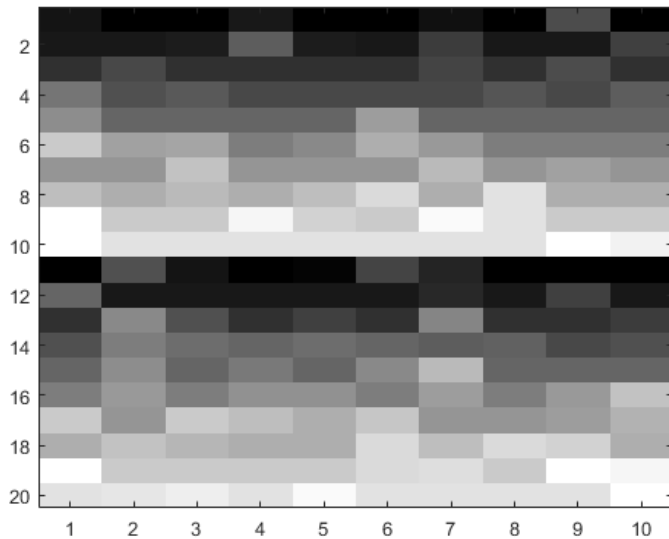
[https://en.wikipedia.org/wiki/Kernel_\(image_processing\)](https://en.wikipedia.org/wiki/Kernel_(image_processing))

I - Basics

Adapt your filters to your problems!

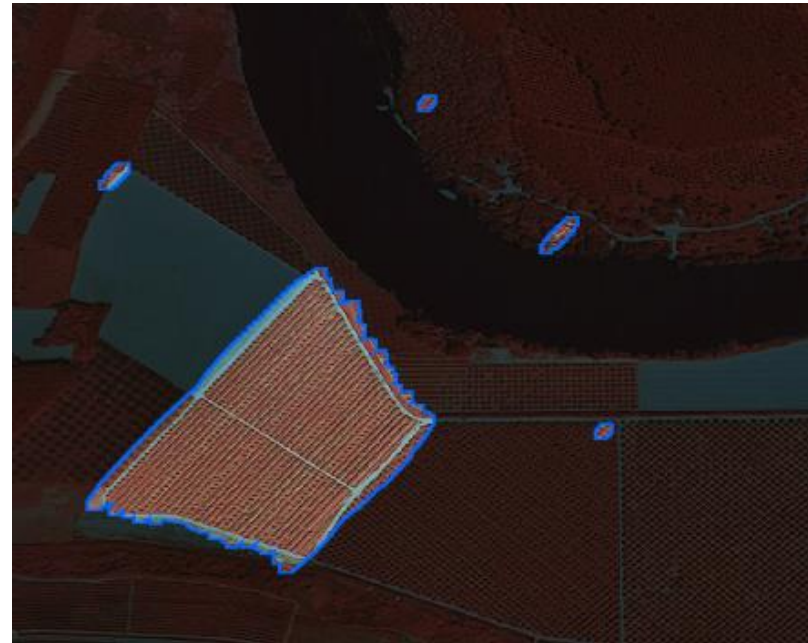
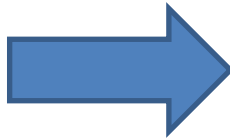
Ex: Smooth noise in image

$$\text{with } H = \frac{1}{n} * [1, \dots, 1]$$



II – Fourier Transform

Frequency analysis



II – Fourier Transform

Reminder: in 1 Dimension

If x is T -periodic

$$x(t) = \frac{1}{2}a_0 + \sum_{k=1}^{\infty} \left(a_k \cos\left(\frac{2k\pi t}{T}\right) + b_k \sin\left(\frac{2k\pi t}{T}\right) \right)$$

$$\left| \begin{array}{l} a_k = \frac{2}{T} \int_0^T x(t) \cos\left(\frac{2k\pi t}{T}\right) dt \\ b_k = \frac{2}{T} \int_0^T x(t) \sin\left(\frac{2k\pi t}{T}\right) dt \end{array} \right.$$


II – Fourier Transform

Reminder: in 1 Dimension

If x is T -periodic

$$x(t) = \frac{1}{2}a_0 + \sum_{k=1}^{\infty} \left(a_k \cos\left(\frac{2k\pi t}{T}\right) + b_k \sin\left(\frac{2k\pi t}{T}\right) \right)$$

$$\langle x(t), \cos\left(\frac{2k\pi t}{T}\right) \rangle$$

$$\left| \begin{array}{l} a_k = \frac{2}{T} \int_0^T x(t) \cos\left(\frac{2k\pi t}{T}\right) dt \\ b_k = \frac{2}{T} \int_0^T x(t) \sin\left(\frac{2k\pi t}{T}\right) dt \end{array} \right.$$


II – Fourier Transform


Reminder: in 1 Dimension

If x is T -periodic

$$x(t) = \frac{1}{2}a_0 + \sum_{k=1}^{\infty} \left(a_k \cos\left(\frac{2k\pi t}{T}\right) + b_k \sin\left(\frac{2k\pi t}{T}\right) \right)$$

$$= \sum_{k=-\infty}^{\infty} r_k \cos\left(\frac{2k\pi t}{T} + \theta_k\right)$$

$$\langle x(t), \cos\left(\frac{2k\pi t}{T}\right) \rangle$$

$$\left| \begin{array}{l} a_k = \frac{2}{T} \int_0^T x(t) \cos\left(\frac{2k\pi t}{T}\right) dt \\ b_k = \frac{2}{T} \int_0^T x(t) \sin\left(\frac{2k\pi t}{T}\right) dt \end{array} \right.$$


II – Fourier Transform

Reminder: in 1 Dimension

If x is T -periodic

$$x(t) = \frac{1}{2}a_0 + \sum_{k=1}^{\infty} \left(a_k \cos\left(\frac{2k\pi t}{T}\right) + b_k \sin\left(\frac{2k\pi t}{T}\right) \right)$$

$$= \sum_{k=-\infty}^{\infty} r_k \cos\left(\frac{2k\pi t}{T} + \theta_k\right)$$

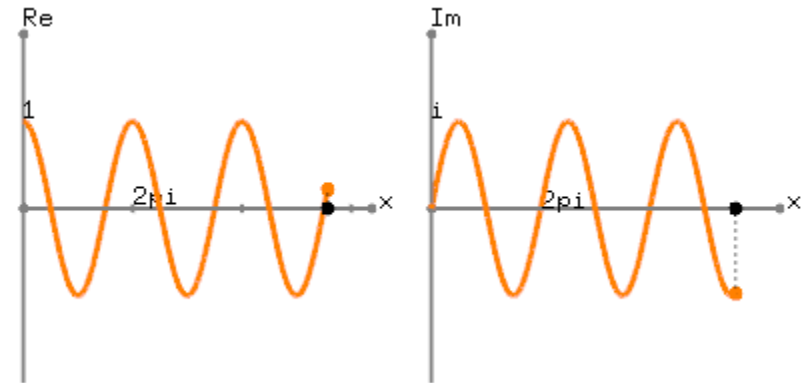
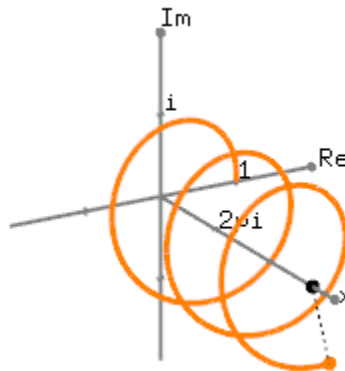
$$= \sum_{k=-\infty}^{\infty} c_k \exp\left(\frac{2ik\pi t}{T}\right)$$

$$\langle x(t), \cos\left(\frac{2k\pi t}{T}\right) \rangle$$

$$a_k = \frac{2}{T} \int_0^T x(t) \cos\left(\frac{2k\pi t}{T}\right) dt$$

$$b_k = \frac{2}{T} \int_0^T x(t) \sin\left(\frac{2k\pi t}{T}\right) dt$$

$$c_k = \frac{1}{T} \int_0^T x(t) \exp\left(\frac{-2ik\pi t}{T}\right) dt$$



II – Fourier Transform

Reminder: in 1 Dimension

If x is T -periodic

$$x(t) = \frac{1}{2}a_0 + \sum_{k=1}^{\infty} \left(a_k \cos\left(\frac{2k\pi t}{T}\right) + b_k \sin\left(\frac{2k\pi t}{T}\right) \right)$$

$$= \sum_{k=-\infty}^{\infty} r_k \cos\left(\frac{2k\pi t}{T} + \theta_k\right)$$

$$= \sum_{k=-\infty}^{\infty} c_k \exp\left(\frac{2ik\pi t}{T}\right)$$

$$\langle x(t), \cos\left(\frac{2k\pi t}{T}\right) \rangle$$

$$\left| \begin{array}{l} a_k = \frac{2}{T} \int_0^T x(t) \cos\left(\frac{2k\pi t}{T}\right) dt \\ b_k = \frac{2}{T} \int_0^T x(t) \sin\left(\frac{2k\pi t}{T}\right) dt \end{array} \right|$$

$$c_k = \frac{1}{T} \int_0^T x(t) \exp\left(\frac{-2i\pi kt}{T}\right) dt$$

$$= \langle x(t), \exp\left(\frac{-2i\pi kt}{T}\right) \rangle$$

Scalar product measures **similarity**

II – Fourier Transform

Reminder: in 1 Dimension

If x is T -periodic

$$x(t) = \frac{1}{2}a_0 + \sum_{k=1}^{\infty} \left(a_k \cos\left(\frac{2k\pi t}{T}\right) + b_k \sin\left(\frac{2k\pi t}{T}\right) \right)$$

$$= \sum_{k=-\infty}^{\infty} r_k \cos\left(\frac{2k\pi t}{T} + \theta_k\right)$$

$$= \sum_{k=-\infty}^{\infty} c_k \exp\left(\frac{2ik\pi t}{T}\right)$$

$$= \sum_{k=-\infty}^{\infty} c_k \exp(2i\pi k f_0 t)$$

$$\langle x(t), \cos\left(\frac{2k\pi t}{T}\right) \rangle$$

$$\left| \begin{array}{l} a_k = \frac{2}{T} \int_0^T x(t) \cos\left(\frac{2k\pi t}{T}\right) dt \\ b_k = \frac{2}{T} \int_0^T x(t) \sin\left(\frac{2k\pi t}{T}\right) dt \end{array} \right.$$

$$c_k = \frac{1}{T} \int_0^T x(t) \exp\left(\frac{-2i\pi k t}{T}\right) dt$$

$$= \langle x(t), \exp\left(\frac{-2i\pi k t}{T}\right) \rangle$$

$$= \langle x(t), \exp(-2i\pi k f_0 t) \rangle$$

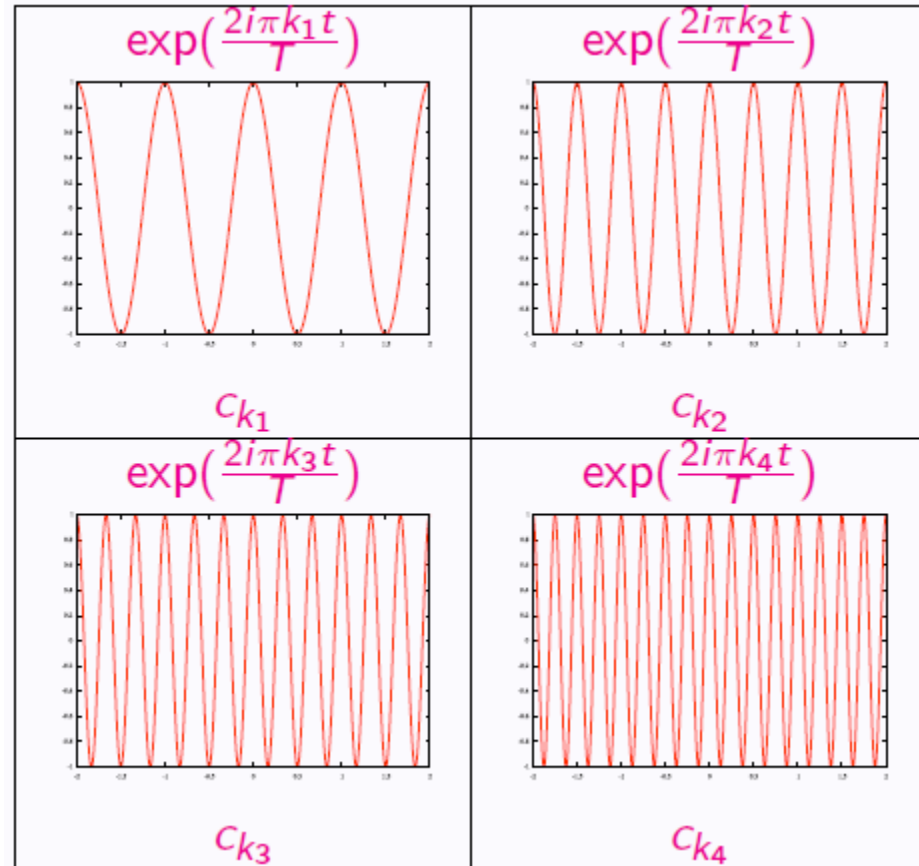
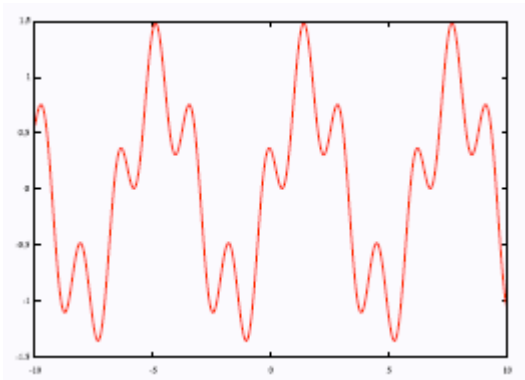
II – Fourier Transform

$$x(t) = \sum_{k=-\infty}^{\infty} c_k \exp\left(\frac{2ik\pi t}{T}\right)$$

$$c_k = \langle x(t), \exp\left(\frac{-2i\pi kt}{T}\right) \rangle$$

$\text{angle}(c_k)$ = phase between $x(t)$ and $\exp(2i\pi k f_0 t)$

$|c_k|$ = intensity of frequency $k f_0$ in $x(t)$



II – Fourier Transform

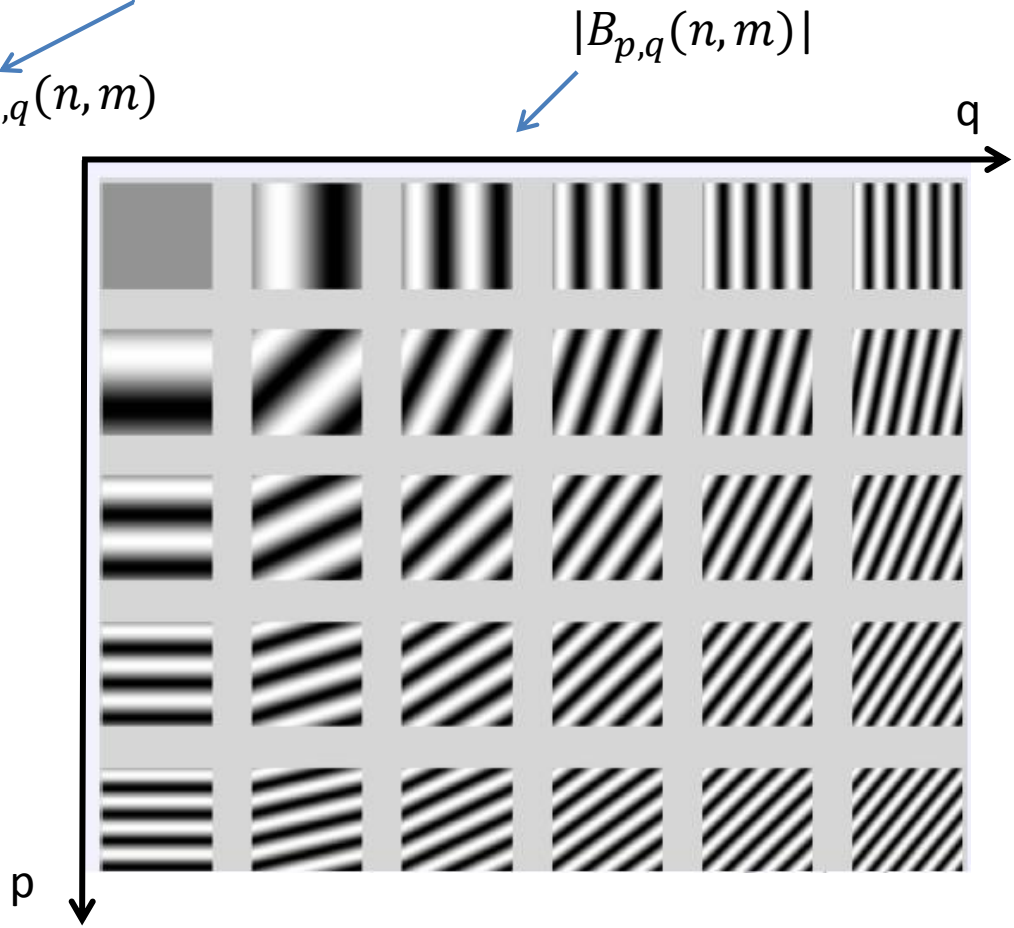
In 2D: $TF: I \rightarrow TF(I) = \hat{I} \in \mathbb{C}$

2 main frequencies: $f_{0_{rows}} = \frac{1}{N}$ and $f_{0_{cols}} = \frac{1}{M}$

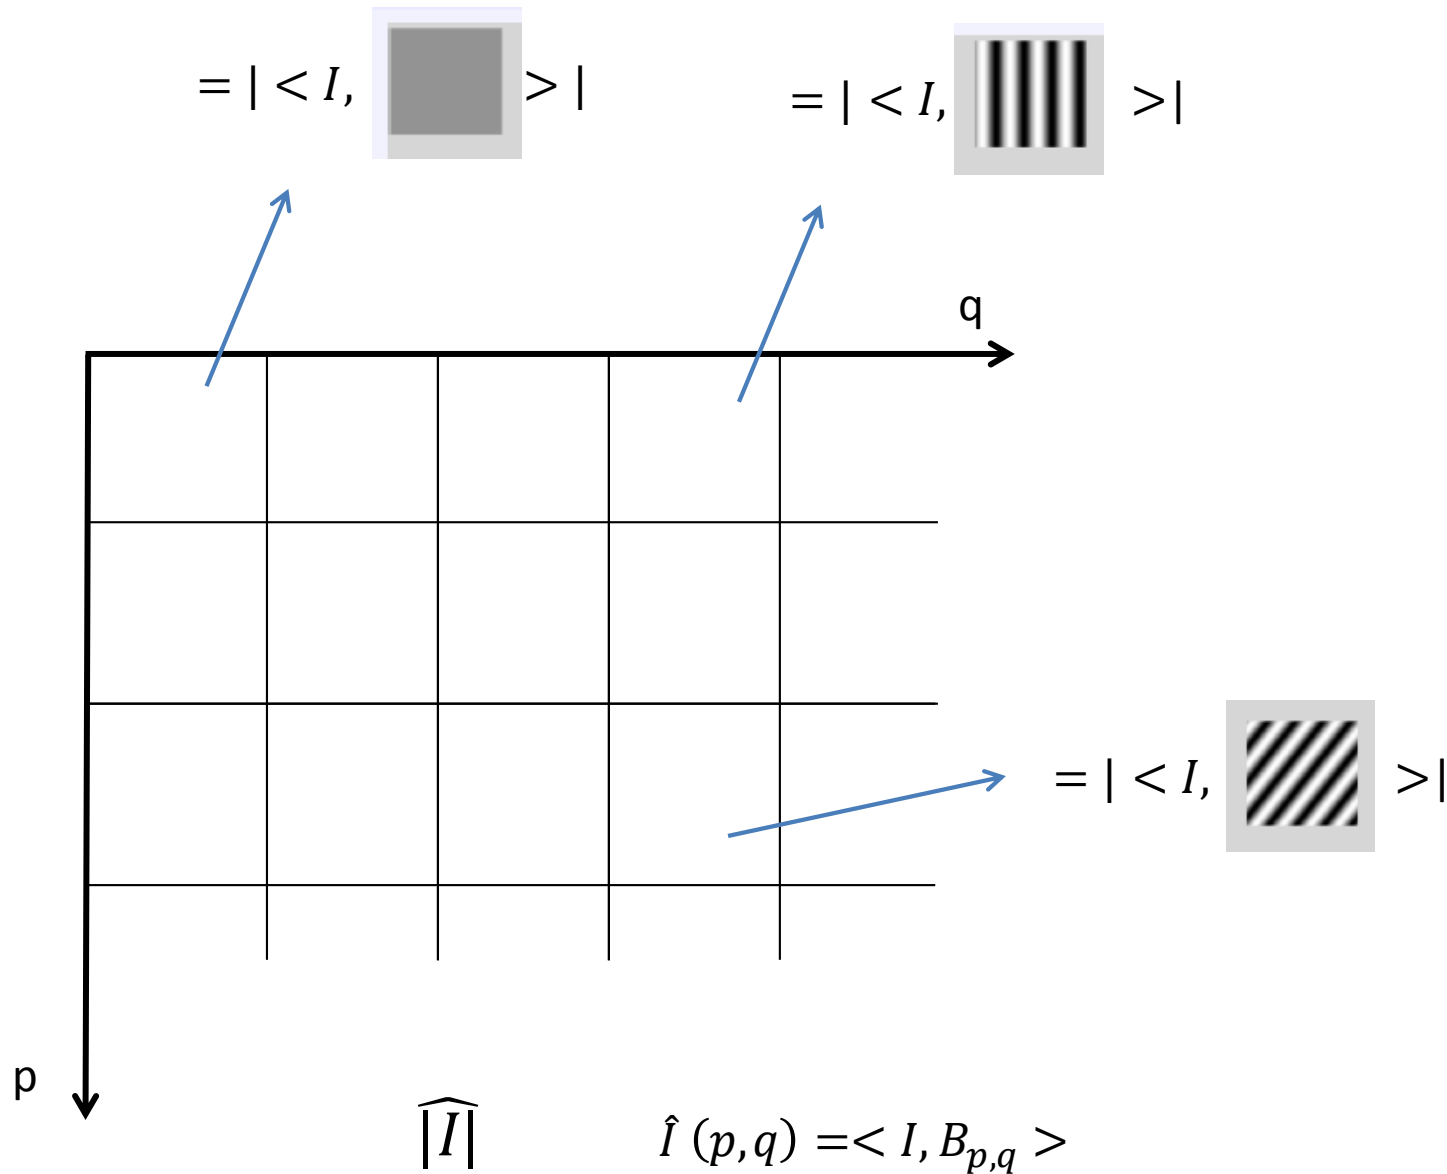
$$\hat{I}(p, q) = \sum_{n=0}^{N-1} \sum_{m=0}^{M-1} I(n, m) \cdot e^{-j\left(\frac{2\pi}{N}\right)pn} e^{-j\left(\frac{2\pi}{M}\right)qm}$$

$$\hat{I}(p, q) = \sum_{n=0}^{N-1} \sum_{m=0}^{M-1} I(n, m) \cdot B_{p,q}(n, m)$$

$$\hat{I}(p, q) = \langle I, B_{p,q} \rangle$$

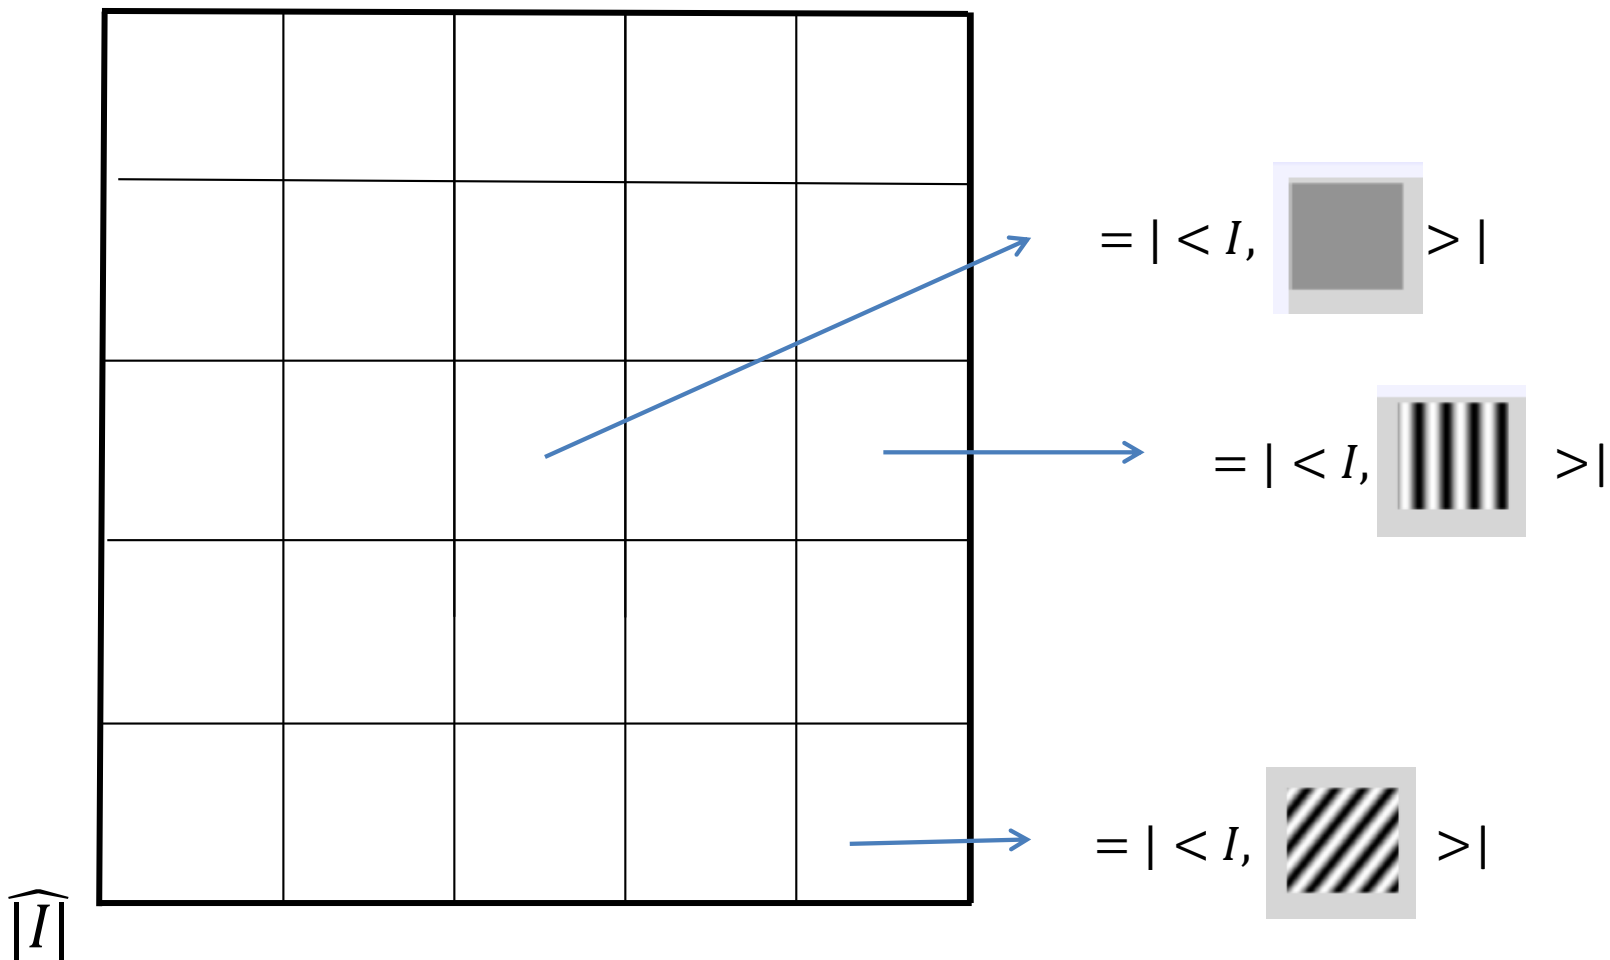


II – Fourier Transform

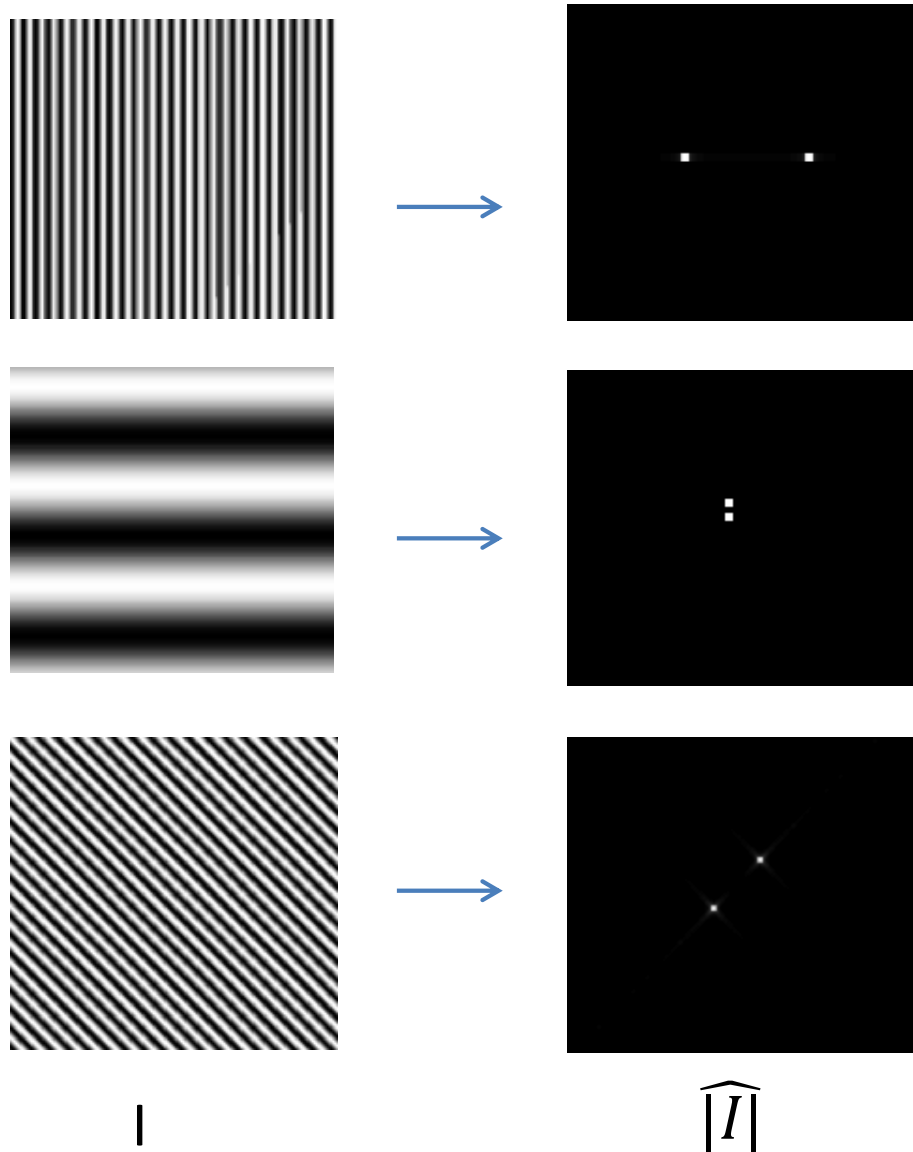


II – Fourier Transform

Warning, conventionally, the low freq are in the middle, and high freq away from the center
Done with function `fftshift`



II – Fourier Transform



Why 2 points?
see:

<https://dsp.stackexchange.com/questions/4825/why-is-the-fft-mirrored>

See Code section 5 intro
(and video 5 intro)

II – Fourier Transform

$$TF: I \rightarrow TF(I) = \hat{I} \in \mathbb{C}$$

$$F(p, q) = \sum_{n=0}^{N-1} \sum_{m=0}^{M-1} I(n, m) \cdot e^{-j(\frac{2\pi}{N})pn} \cdot e^{-j(\frac{2\pi}{M})qm}$$

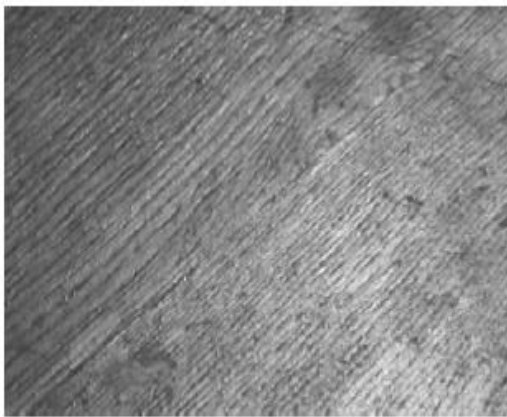
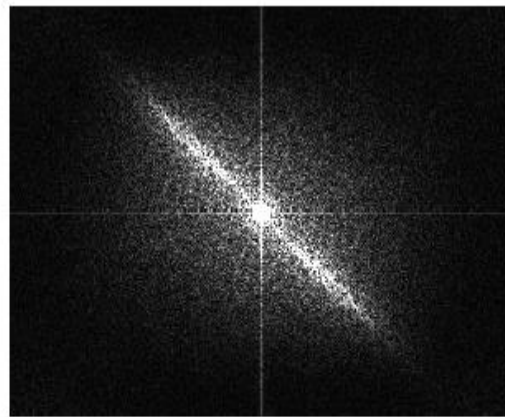
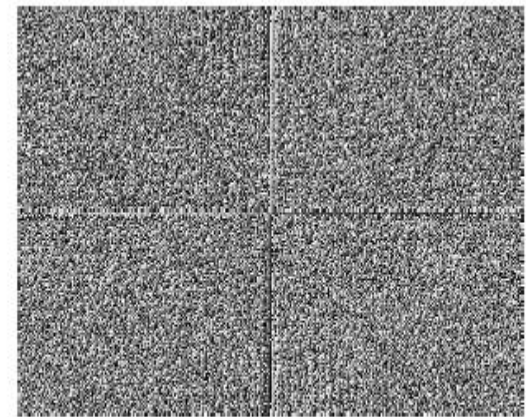


Image u



Module $|\hat{u}|$

\approx Main directions/Patterns



Phase $\arg(\hat{u})$

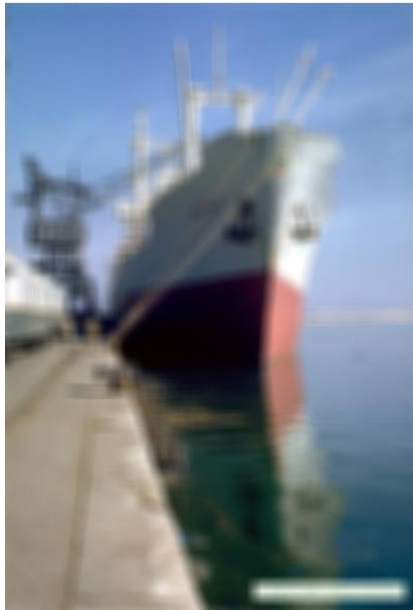
\approx Texture
spatial positionning

- $I_{fft} = fftshift(fft2(I))$
- `imagesc(log(abs(I_fft)))` % Log in order to see more details
- $I = ifft2(fftshift(I_{fft}))$

Important : $TF(H * I) = TF(H) \times TF(I)$

See Code section 5 (and video 5)

III – Deblurring



$$I_b = H * I + N$$



Filtering: N is attenuated

$$I_{b_{filtered}} = H * I + n$$



$$I_{deblurred} = \frac{I_{b_{filtered}}}{H_{est}}$$



Problems:

- How to efficiently remove N (filtering)?
- How to correctly model H ?

Annexe – TF for non periodic signals

If x is not T -periodic, replace $\frac{k}{T} = kf_0$ with f , \sum_k with \int_f , and c_k with $\hat{X}(f)$

$$c_k = \langle x(t), \exp(-2i\pi kf_0 t) \rangle \longrightarrow \hat{X}(f) = \langle x(t), \exp(-2i\pi f t) \rangle$$

$$x(t) = \sum_{k=-\infty}^{\infty} c_k \exp(2i\pi kf_0 t) \longrightarrow x(t) = \int_{f=-\infty}^{\infty} \hat{X}(f) \exp(2i\pi f t) df$$

$$TF: x \rightarrow TF(x) = \hat{X} \in \mathbb{C}$$

$|\hat{X}(f)|$ = intensity of frequency f in $x(t)$

$\angle(\hat{X}(f))$ = phase between $x(t)$ and $\exp(2i\pi f t)$